

# CSSS 510: Lab 2

## Introduction to Maximum Likelihood Estimation

2017-10-6

### 0. Agenda

1. Simulating heteroskedastic normal data
2. Fitting a model using the simulated data
3. Calculating predicted values
4. Fitting the heteroskedastic normal model using ML
5. Simulating predicted values and confidence intervals

### 1. Simulating heteroskedastic normal data

Stochastic component:

$$y \sim N(\mu_i, \sigma_i^2)$$

Systematic components:

$$\mu = \mathbf{x}_i \boldsymbol{\beta}$$

$$\sigma_i^2 = \exp(\mathbf{z}_i \boldsymbol{\gamma})$$

- Set the number of observations to 1500 ( $n$ )
- Set a parameter vector for the mean (assume 2 covariates plus the constant) ( $\boldsymbol{\beta}$ )
- Set a parameter vector for the variance (assume heteroskedasticity) ( $\boldsymbol{\gamma}$ )
- Generate the constant and the covariates, length 1500 for each (draw from a uniform distribution) ( $\mathbf{x}_i$ ,  $\mathbf{z}_i$ )
- Create the systematic component for the mean ( $\mathbf{x}_i \boldsymbol{\beta}$ )
- Create the systematic component for the variance (the same covariates affect mu and sigma)  $\exp(\mathbf{z}_i \boldsymbol{\gamma})$
- Generate the response variable ( $y_i$ )
- Save the data to a data frame
- Plot the data

```
rm(list=ls()) # Clear memory
set.seed(123456) # To reproduce random numbers
library(MASS) # Load packages
library(simcf)
```

```

n <- 1500 # Generate 1500 observations

beta <- c(0, 5, 15) # Set a parameter vector for the mean
# One for constant, one for covariate 1, one for covariate 2.

gamma <- c(1, 0, 3) # Set a parameter vector for the variance
# Note that gamma estimate for covariate 2 is set to be 3, creating heteroskedasticity

w0 <- rep(1, n) # Create the constant and covariates, length of each vector is 1500
w1 <- runif(n)
w2 <- runif(n)

x <- cbind(w0, w1, w2) # Create a matrix of the covariates

mu <- x%*%beta # Create the systematic component for the mean

z <- x # i.e., same covariates affect mu and sigma
sigma2 <- exp(x%*%gamma) # Create the systematic component for the variance

# z is 1500 by 3 matrix, gamma is 3 by 1 matrix
# ith row of sigma 2 thus equals exp(1+0+w2_i*3). i.e., it is a function of w2

y <- mu + rnorm(n)*sqrt(sigma2) # Create the response variable

data <- cbind(y,w1,w2) # Save the data to a data frame
data <- as.data.frame(data)
names(data) <- c("y","w1","w2")

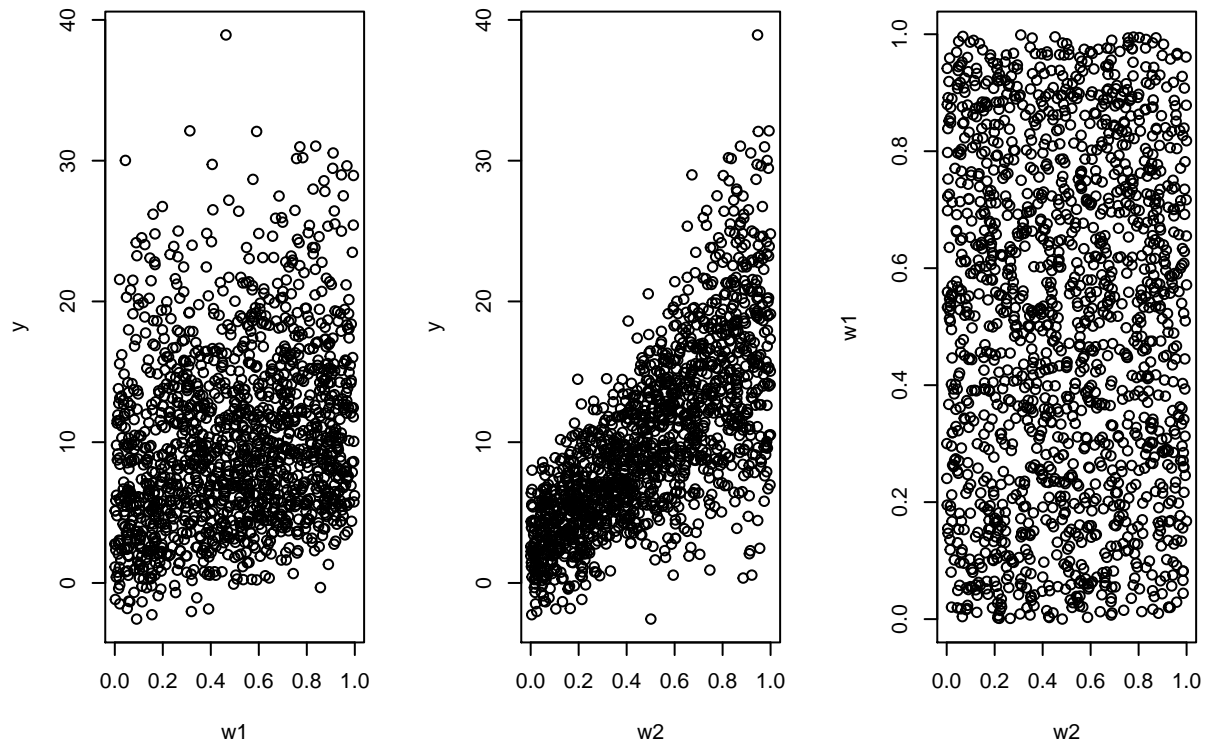
par(mfrow=c(1,3)) #Plot the data

#pdf("YvsW1.pdf")
plot(y=y,x=w1)
#dev.off()

#pdf("YvsW2.pdf")
plot(y=y,x=w2)
#dev.off()

#pdf("W1vsW2.pdf")
plot(y=w1,x=w2)

```



```
#dev.off()
```

## 2. Fitting a model using the simulated data

- Assume we don't know the true value of the parameters and fit a model using least squares (use the `lm()` function and regress the response variable on the two covariates)
- Calculate and print the AIC

```
ls.result <- lm(y ~ w1 + w2) #Fit a linear model using the simulated data
print(summary(ls.result))
```

```
##
## Call:
## lm(formula = y ~ w1 + w2)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-16.5710	-2.3157	-0.0219	2.2124	21.9345

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-0.04867	0.28221	-0.172	0.863
w1	4.78421	0.37699	12.690	<2e-16 ***
w2	15.68283	0.37112	42.258	<2e-16 ***

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
##
## Residual standard error: 4.17 on 1497 degrees of freedom
## Multiple R-squared:  0.5678, Adjusted R-squared:  0.5672
## F-statistic: 983.3 on 2 and 1497 DF,  p-value: < 2.2e-16

# Calculate and print the AIC
ls.aic <- AIC(ls.result)
print(ls.aic)

## [1] 8545.903
```

### 3. Calculating predicted values

#### Scenario 1: Vary covariate 1

- Create a data frame with a set of hypothetical scenarios for covariate 1 while keeping covariate 2 at its mean
- Calculate the predicted values using the `predict()` function
- Plot the predicted values

```
# Calculate predicted values using predict()

# Start by calculating  $P(Y|w1)$  for different  $w1$  values
w1range <- seq(0:20)/20 # Set as necessary

# Set up a dataframe with the hypothetical scenarios (varied  $w1$ , all else equal)
baseline <- c(mean(w1), mean(w2)) # Set as necessary
xhypo <- matrix(baseline, nrow=length(w1range), ncol=2, byrow= TRUE) # Set ncol to # of  $x$ 's
# same as: xhypo <- matrix(rep(baseline,21), nrow=length(w1range), ncol=2, byrow= TRUE)

xhypo <- as.data.frame(xhypo)
names(xhypo) <- c("w1", "w2")
xhypo[,1] <- w1range
# Scenarios: Changing values in the first column, keeping second column values at the mean of  $w2$ .
xhypo
```

```
##      w1      w2
## 1  0.05 0.4912698
## 2  0.10 0.4912698
## 3  0.15 0.4912698
## 4  0.20 0.4912698
## 5  0.25 0.4912698
## 6  0.30 0.4912698
## 7  0.35 0.4912698
## 8  0.40 0.4912698
## 9  0.45 0.4912698
## 10 0.50 0.4912698
## 11 0.55 0.4912698
## 12 0.60 0.4912698
## 13 0.65 0.4912698
```

```
## 14 0.70 0.4912698
## 15 0.75 0.4912698
## 16 0.80 0.4912698
## 17 0.85 0.4912698
## 18 0.90 0.4912698
## 19 0.95 0.4912698
## 20 1.00 0.4912698
## 21 1.05 0.4912698


# Calculate Predicted Y using predict()
simls.w1 <- predict(ls.result, newdata = xhypo, interval = "prediction", level = 0.95)
simls.w1
```

##	fit	lwr	upr
## 1	7.895048	-0.29514691	16.08524
## 2	8.134259	-0.05450529	16.32302
## 3	8.373470	0.18596961	16.56097
## 4	8.612680	0.42627769	16.79908
## 5	8.851891	0.66641891	17.03736
## 6	9.091101	0.90639320	17.27581
## 7	9.330312	1.14620051	17.51442
## 8	9.569523	1.38584081	17.75320
## 9	9.808733	1.62531406	17.99215
## 10	10.047944	1.86462027	18.23127
## 11	10.287154	2.10375941	18.47055
## 12	10.526365	2.34273150	18.71000
## 13	10.765575	2.58153654	18.94961
## 14	11.004786	2.82017457	19.18940
## 15	11.243997	3.05864562	19.42935
## 16	11.483207	3.29694973	19.66946
## 17	11.722418	3.53508696	19.90975
## 18	11.961628	3.77305737	20.15020
## 19	12.200839	4.01086104	20.39082
## 20	12.440050	4.24849806	20.63160
## 21	12.679260	4.48596852	20.87255

```
# Plot them
yplot <- simls.w1
xplot <- cbind(w1range,w1range,w1range) # need to have the same dimension [21,3]
#pdf("homoYusw1.pdf")

# matplotlib is useful for plotting matrices of data

matplot(y=yplot,
        x=xplot,
        type="l",
        lty=c("solid","dashed","dashed"),
        col=c("black"),
        xlab = "w1",
        ylab = "Predicted Y")
```



```
#dev.off()
```

## Scenario 2: Vary covariate 2

- Create a data frame with a set of hypothetical scenarios for covariate 2 while keeping covariate 1 at its mean
- Calculate the predicted values using the `predict()` function
- Plot the predicted values

```
# Calculate predicted values using predict()


# Start by calculating P(Y|w1) for different w1 values
w2range <- seq(0:20)/20 # Set as necessary

# Set up a dataframe with the hypothetical scenarios (varied w1, all else equal)
baseline <- c(mean(w1), mean(w2)) # Set as necessary
xhypo <- matrix(baseline, nrow=length(w2range), ncol=2, byrow= TRUE) # Set ncol to # of x's
xhypo <- as.data.frame(xhypo)
names(xhypo) <- c("w1", "w2") # Set by user
xhypo[,2] <- w2range # Change as necessary

# Calculate Predicted Y using predict()
simls.w2 <- predict(ls.result, newdata = xhypo, interval = "prediction", level = 0.95)

# Plot them
yplot <- simls.w2
xplot <- cbind(w2range, w2range, w2range)
#pdf("homoYusw2.pdf")
```

```
matplot(y=yplot,
        x=xplot,
        type="l",
        lty=c("solid","dashed","dashed"),
        col=c("black"),
        xlab = "w1",
        ylab = "Predicted Y")
```



```
#dev.off()
```

## 4. Fitting the heteroskedastic normal model using ML

- Create the input matrices (the two covariates)
- Write a likelihood function for the heteroskedastic normal model
- Find the MLEs using the `optim()` function
- Extract the point estimates
- Compute the standard errors
- Compare with the least squares estimates
- Find the log likelihood at its maximum
- Compute the AIC
- Simulate the results by drawing from the model's predictive distribution
- Separate the simulated betas from the simulated gammas

Recall from lecture:

$$\begin{aligned}\mathcal{L}(\boldsymbol{\mu}, \sigma^2 | \mathbf{y}) &\propto P(\mathbf{y} | \boldsymbol{\mu}, \sigma^2) \\ \mathcal{L}(\boldsymbol{\mu}, \sigma^2 | \mathbf{y}) &= k(\mathbf{y}) P(\mathbf{y} | \boldsymbol{\mu}, \sigma^2) \\ \mathcal{L}(\boldsymbol{\mu}, \sigma^2 | \mathbf{y}) &= k(\mathbf{y}) \prod_{i=1}^n (2\pi\sigma^2)^{-1/2} \exp\left(\frac{-(y_i - \mu_i)^2}{2\sigma^2}\right) \\ &\dots \\ \mathcal{L}(\boldsymbol{\beta}, \sigma^2 | \mathbf{y}) &= -\frac{1}{2} \sum_{i=1}^n \log \sigma^2 - \frac{1}{2} \sum_{i=1}^n \frac{(y_i - \mathbf{x}_i \boldsymbol{\beta})^2}{\sigma^2} \\ \mathcal{L}(\boldsymbol{\beta}, \boldsymbol{\gamma} | \mathbf{y}) &= -\frac{1}{2} \sum_{i=1}^n z_i \gamma - \frac{1}{2} \sum_{i=1}^n \frac{(y_i - \mathbf{x}_i \boldsymbol{\beta})^2}{\exp(z_i \gamma)}\end{aligned}$$

```
# A likelihood function for ML heteroskedastic Normal
llk.hetnormlin <- function(param,y,x,z) {
  x <- as.matrix(x)      #x as a matrix
  z <- as.matrix(z)      #z as a matrix
  os <- rep(1,nrow(x))  #1 for the intercept
  x <- cbind(os,x)       #combine
  z <- cbind(os,z)
  b <- param[ 1 : ncol(x) ] # i.e., the first three spaces in the param vector
  g <- param[ (ncol(x)+1) : (ncol(x) + ncol(z)) ] # i.e., the three remaining spaces
  xb <- x%*%b # systematic components for the mean
  s2 <- exp(z%*%g) # systematic components for the variance

  sum(0.5*(log(s2)+(y-xb)^2/s2)) # "optim" command minimizes a function by default.
  # Minimalization of -lnL is the same as maximization of lnL, so we will put -lnL(param/y) here

  #-sum(0.5*(log(s2)+(y-xb)^2/s2))
  # Alternativly, you can use lnL(param/y) and set optim to be a maximizer

}

# Create input matrices
xcovariates <- cbind(w1,w2)
zcovariates <- cbind(w1,w2)

# initial guesses of beta0, beta1, ..., gamma0, gamma1, ...
# we need one entry per parameter, in order!
stval <- c(0,0,0,0,0,0) # also include beta and gamma estimates for constants

help(optim)

# Run ML, get the output we need
hetnorm.result <- optim(stval,llk.hetnormlin,method="BFGS",hessian=T,y=y,x=xcovariates,z=zcovariates)
# by default, calls minimizer procedure.
# you can make optim a maximizer by adding control=list(fnscale=-1)

pe <- hetnorm.result$par # point estimates
pe
```



```
## [1] -0.1672393  5.0074031 15.6981262  0.9103126  0.2238205  2.9937686

vc <- solve(hetnorm.result$hessian) # 6x6 var-cov matrix (allows to compute standard errors)
round(vc,5)

##          [,1]      [,2]      [,3]      [,4]      [,5]      [,6]
## [1,]  0.02909 -0.03265 -0.02810  0.00059 -0.00087 -0.00032
## [2,] -0.03265  0.06787 -0.00025 -0.00091  0.00099  0.00084
## [3,] -0.02810 -0.00025  0.10331 -0.00056  0.00143 -0.00033
## [4,]  0.00059 -0.00091 -0.00056  0.00920 -0.00832 -0.00749
## [5,] -0.00087  0.00099  0.00143 -0.00832  0.01693 -0.00042
## [6,] -0.00032  0.00084 -0.00033 -0.00749 -0.00042  0.01568

se <- sqrt(diag(vc)) # standard errors
# the ML standard errors are the square roots of the diagonal of the Hessian
# or inverse of the matrix of second derivatives
se

## [1] 0.17056695 0.26050990 0.32141276 0.09594005 0.13010517 0.12523484

mle.result<-round(cbind(pe[1:3], se[1:3]),2) # see pe and se
colnames(mle.result)<-c("Estimate", "Std.Error")
rownames(mle.result)<-c("(Intercept)", "w1", "w2" )
mle.result

##              Estimate Std. Error
## (Intercept)   -0.17      0.17
## w1              5.01      0.26
## w2             15.70      0.32

round(summary(ls.result)$coefficients[,c(1,2)],2) #compare with the ls result

##              Estimate Std. Error
## (Intercept)   -0.05      0.28
## w1              4.78      0.38
## w2             15.68      0.37

ll <- -hetnorm.result$value
# likelihood at maximum, no need to have a negative sign if you set optime to be a maximizer.
ll

## [1] -2620.334

#The AIC is the deviance or -2*ll at its max plus 2*number of parameters or the dimension
hetnorm.aic <- 2*length(stval) - 2*ll
# first component to penalizing the number of parameters
# (i.e., the loss of degree of freedom). Lower aic is better

print(hetnorm.aic)

## [1] 5252.668

# remember AIC from LS fit?
print(ls.aic)

## [1] 8545.903
```

## 5. Simulating predicted values and confidence intervals

### Scenario 1: Vary covariate 1

- Create a data frame with a set of hypothetical scenarios for covariate 1 while keeping covariate 2 at its mean
- Simulate the predicted values and the confidence intervals using simcf
- Plot the results

```
# Simulate results by drawing from the model predictive distribution
sims <- 10000
simparam <- mvrnorm(sims,pe,vc) # draw parameters store them in 10000x6 matrix.
# We assume that parameter estimates are distributed according to a multivariate normal
# distribution with population mean pe and population variance-covariance matrix vc.

# Separate into the simulated betas and simulated gammas
simbetas <- simparam[,1:(ncol(xcovariates)+1)]
# first three columns store simulated beta coefficients
simgammas <- simparam[(ncol(simbetas)+1):ncol(simparam)]
# then simulated gamma coefficients

# Put our models in "formula" form
model <- (y ~ w1 + w2)
varmodel <- (y ~ w1 + w2)

# Scenario 1: Vary w1

# Start by calculating P(Y|w1) for different w1 values
w1range <- seq(0:20)/20

# Set up a matrix with the hypothetical scenarios (varied w1, all else equal)
xhypo <- cfMake(model, data, nscen = length(w1range)) # creating a set of scenarios

for (i in 1:length(w1range)) {
  xhypo <- cfChange(xhypo, "w1", x=w1range[i], scen=i)
  # change the values of the variables of your interest, set others at the mean
}

zhypo <- cfMake(varmodel, data, nscen = length(w1range))

for (i in 1:length(w1range)) {
  zhypo <- cfChange(zhypo, "w1", x=w1range[i], scen=i)
}


# Simulate the predicted Y's and CI's
simres.w1 <- hetnormsimpv(xhypo,simbetas,
                          zhypo,simgammas,
                          ci=0.95,
                          constant=1,varconstant=1)
```

```

#simy<-rnorm(sims)*sqrt(simsigma2)+ simmu

# Plot them
yplot <- cbind(simres.w1$pe, simres.w1$lower, simres.w1$upper)
xplot <- cbind(w1range,w1range,w1range)
#pdf("heteroYvsW1.pdf")
matplot(y=yplot,
        x=xplot,
        type="l",
        lty=c("solid","dashed","dashed"),
        col=c("black"),
        xlab = "w1",
        ylab = "Predicted Y")

```



```

#dev.off()

```

## Scenario 2: Vary covariate 2

- Create a data frame with a set of hypothetical scenarios for covariate 2 while keeping covariate 1 at its mean
- Simulate the predicted values and the confidence intervals using simcf
- Plot the results

```

# Start by calculating P(Y/w2) for different w1 values
w2range <- seq(0:20)/20

# Set up a matrix with the hypothetical scenarios (varied w1, all else equal)
xhypo <- cfMake(model, data, nscen = length(w2range))

```

```


for (i in 1:length(w2range)) {
  xhypo <- cfChange(xhypo, "w2", x=w2range[i], scen=i)
}

zhypo <- cfMake(varmodel, data, nscen = length(w2range))
for (i in 1:length(w2range)) {
  zhypo <- cfChange(zhypo, "w2", x=w2range[i], scen=i)
}

# Simulate the predicted Y's and CI's
simres.w2 <- hetnormsimpv(xhypo,simbetas,
                        zhypo,simgammas,
                        ci=0.95,
                        constant=1,varconstant=1)

# Plot them
yplot <- cbind(simres.w2$pe, simres.w2$lower, simres.w2$upper)
xplot <- cbind(w1range,w1range,w1range)
#pdf("heteroYvsW2.pdf")
matplot(y=yplot,
        x=xplot,
        type="l",
        lty=c("solid","dashed","dashed"),
        col=c("black"),
        xlab = "w2",
        ylab = "Predicted Y")

```



```

# since we set the
#dev.off()

```