

CSSS 510: Lab 5

Ordered Probit

2017-11-10

0. Agenda

1. Deriving a likelihood function for the ordered probit model
2. Fitting an ordered probit model using `optim()` and `glm()`
3. Interpreting the results
4. Simulating predicted values and confidence intervals

1. Deriving a likelihood function for ordered probit

Recall from lecture the ordered probit model:

$$\Pr(y_i = j | \mathbf{x}_i) = \int_{\tau_{j-1}}^{\tau_j} \text{Normal}(\mathbf{x}_i \beta, 1) d\mathbf{x}_i \beta$$

How does this model differ from the others we've covered so far in the course?


We are saying that the probability that y_i is in category j is equal to the CDF of the standard normal distribution evaluated at $\mathbf{x}_i \beta$ between cutpoints τ_j and τ_{j-1}

1. Deriving a likelihood function for ordered probit

If we merely plot the observed values of y then we would obtain something like the following.

But, this runs into the same problems we found when using linear regression to fit a dichotomous outcome variable.

Panel B: Regression of Observed y



1. Deriving a likelihood function for ordered probit


Instead, we refer back to the latent variable framework, which produces a similar relationship between our covariates and y^* but now we have more than two categories.

In a four category model, we have the following:

$$y_i = \begin{cases} 1 \Rightarrow \text{SD}, & \text{if } \tau_0 = -\infty \leq y_i^* < \tau_1 \\ 2 \Rightarrow \text{D}, & \text{if } \tau_1 \leq y_i^* < \tau_2 \\ 3 \Rightarrow \text{A}, & \text{if } \tau_2 \leq y_i^* < \tau_3 \\ 4 \Rightarrow \text{SA}, & \text{if } \tau_3 \leq y_i^* < \tau_4 = \infty \end{cases}$$

$$y^* = \mathbf{x}_i \boldsymbol{\beta} + \epsilon$$

Panel A: Regression of Latent y^*



1. Deriving a likelihood function for ordered probit

Recall that in logit, we assume the errors of the latent variable follow a standard logistic distribution.

For probit and ordered probit, we assume the errors follow a standard normal distribution.




Figure 5.2. Distribution of y^* Given x for the Ordered Regression Model

$$\begin{aligned}\Pr(y_i = 1 | \mathbf{x}_i) &= \Pr(\tau_0 \leq y^* < \tau_1 | \mathbf{x}_i) \\ &= \Pr(\tau_0 \leq \mathbf{x}_i \beta + \epsilon_i < \tau_1 | \mathbf{x}_i) \\ &= \Pr(\tau_0 - \mathbf{x}_i \beta \leq \epsilon_i < \tau_1 - \mathbf{x}_i \beta | \mathbf{x}_i) \\ &= \Pr(\epsilon_i < \tau_1 - \mathbf{x}_i \beta | \mathbf{x}_i) - \Pr(\epsilon_i \leq \tau_0 - \mathbf{x}_i \beta | \mathbf{x}_i) \\ &= F(\tau_1 - \mathbf{x}_i \beta) - F(\tau_0 - \mathbf{x}_i \beta) \\ \Pr(y_i = m | \mathbf{x}_i) &= F(\tau_m - \mathbf{x}_i \beta) - F(\tau_{m-1} - \mathbf{x}_i \beta)\end{aligned}$$

1. Deriving a likelihood function for ordered probit

For the ordered probit this becomes




Figure 5.2. Distribution of y^* Given x for the Ordered Regression Model

$$\Pr(y_i = 1|\mathbf{x}_i) = \Phi(\tau_1 - \alpha - \mathbf{x}_i\beta)$$

$$\Pr(y_i = 2|\mathbf{x}_i) = \Phi(\tau_2 - \alpha - \mathbf{x}_i\beta) - \Phi(\tau_1 - \alpha - \mathbf{x}_i\beta)$$

$$\Pr(y_i = 3|\mathbf{x}_i) = \Phi(\tau_3 - \alpha - \mathbf{x}_i\beta) - \Phi(\tau_2 - \alpha - \mathbf{x}_i\beta)$$

$$\Pr(y_i = 4|\mathbf{x}_i) = 1 - \Phi(\tau_3 - \alpha - \mathbf{x}_i\beta)$$

1. Deriving a likelihood function for ordered probit

To identify the model, we commonly make one of two assumptions:

1. Assume that $\tau_1 = 0$. This is also the identifying assumption of logit and probit. `optim()` uses this.
2. Assume that $\alpha = 0$. `polr()` uses this.

The likelihood function for ordered probit finds the β and τ that make the observed data most likely.

$$\mathcal{L}(\beta, \tau | \mathbf{y}, \mathbf{X}) = \prod_{i=1}^n \left\{ \prod_{j=1}^m [\Phi(\tau_j | \mathbf{x}_i \beta, 1) - \Phi(\tau_{j-1} | \mathbf{x}_i \beta, 1)]^{y_{ij}} \right\}$$

$$\mathcal{L}(\beta, \tau | \mathbf{y}, \mathbf{X}) = \sum_{i=1}^n \sum_{j=1}^m y_{ij} \log[\Phi(\tau_j | \mathbf{x}_i \beta, 1) - \Phi(\tau_{j-1} | \mathbf{x}_i \beta, 1)]$$

We estimate β and τ that maximizes the likelihood that y_{ij} falls into category j . All other categories $\neq j$ are irrelevant.

2. Fitting an ordered probit model using optim() and glm()

```
rm(list=ls())

## Likelihood for 4 category ordered probit
llk.oprobit4 <- function(param, x, y) {
  # preliminaries
  os <- rep(1, nrow(x))
  x <- cbind(os, x)
  b <- param[1:ncol(x)]
  t2 <- param[(ncol(x)+1)]
  t3 <- param[(ncol(x)+2)]

  # probabilities and penalty function
  xb <- x%*%b
  p1 <- log(pnorm(-xb))
  if (t2<=0) p2 <- -(abs(t2)*10000)      # penalty function to keep t2>0
  else p2 <- log(pnorm(t2-xb)-pnorm(-xb))
  if (t3<=t2) p3 <- -((t2-t3)*10000)    # penalty to keep t3>t2
  else p3 <- log(pnorm(t3-xb)-pnorm(t2-xb))
  p4 <- log(1-pnorm(t3-xb))

  # -1 * log likelihood (optim is a minimizer)
  -sum(cbind(y==1,y==2,y==3,y==4) * cbind(p1,p2,p3,p4))
}
```

2. Fitting an ordered probit model using `optim()` and `glm()`

```
## Load libraries
library(MASS)
library(simcf)
library(tile)
```

```
## Loading required package: grid
```

```
library(RColorBrewer)
```

```
## Nice colors
brewer <- brewer.pal(9, "Set1")
red <- brewer[1]
blue <- brewer[2]
green <- brewer[3]
purple <- brewer[4]
orange <- brewer[5]
nicegray <- "gray45"
```

```
## Load data
workmom <- read.csv("ordwarm2.csv", header=TRUE, sep=",")
workmom77 <- workmom[workmom$yr89==0, ]
workmom89 <- workmom[workmom$yr89==1, ]
```

2. Fitting an ordered probit model using optim() and glm()

```
## Data from 1977, 1989 GSS: Attitudes towards working mothers
y <- workmom77$warm      # Mother can have warm feelings towards child?

x <- cbind(workmom77$male, workmom77$white, workmom77$age,
            workmom77$ed, workmom77$prst)
## male respondent; white resp; age of resp;
## years of education of respondent;
## prestige of respondent's occupation (% considering prestigious)

# Model specification (for polr, simcf)
model <- warm ~ male + white + age + ed + prst

# Use optim directly to get MLE
ls.result <- lm(model, data=workmom77)    # use ls estimates as starting values
stval <- c(coef(ls.result),1,2)           # initial guesses
oprobit.res77 <- optim(stval, llk.oprobit4, method="BFGS", x=x, y=y, hessian=T)
pe77 <- oprobit.res77$par                  # point estimates
vc77 <- solve(oprobit.res77$hessian)       # var-cov matrix
se77 <- sqrt(diag(vc77))                  # standard errors
ll77 <- -oprobit.res77$value               # likelihood at maximum
```

2. Fitting an ordered probit model using `optim()` and `glm()`

```
pe77
```

```
## (Intercept)          male          white          age          ed
## 1.321689587 -0.397228916 -0.238335818 -0.011703265 0.044989779
##          prst
## 0.002116029 1.016421525 2.077970893
```

```
se77
```

```
## (Intercept)          male          white          age          ed          prst
## 0.179052337 0.058429471 0.091086552 0.001915842 0.012076026 0.002586886
##
## 0.041094751 0.054430991
```

```
ll77
```

```
## [1] -1758.824
```

4. Simulating predicted values and confidence intervals

```
# Use MASS::polr to do ordered probit
workmom77$warmf <- factor(workmom77$warm, labels=c("Strongly Disagree",
                                                    "Disagree",
                                                    "Agree",
                                                    "Strongly Agree"))

glm.res77 <- polr(warmf ~ male + white + age + ed + prst, data=workmom77,
                 method="probit", na.action=na.omit)

# Simulate parameters from predictive distributions
sims <- 10000
simbetas <- mvrnorm(sims, pe77, vc77) # draw parameters, using MASS::mvrnorm

# Create example counterfactuals
xhyp <- cfMake(model, workmom77, nscen=10)

xhyp <- cfName(xhyp, "Male", scen=1)
xhyp <- cfChange(xhyp, "male", x=1, xpre=0, scen=1)

xhyp <- cfName(xhyp, "Female", scen=2)
xhyp <- cfChange(xhyp, "male", x=0, xpre=1, scen=2)

xhyp <- cfName(xhyp, "Nonwhite", scen=3)
xhyp <- cfChange(xhyp, "white", x=0, xpre=1, scen=3)

xhyp <- cfName(xhyp, "White", scen=4)
xhyp <- cfChange(xhyp, "white", x=1, xpre=0, scen=4)
```

4. Simulating predicted values and confidence intervals

```
xhyp <- cfName(xhyp, "Age + 1sd = 61", scen=5)
xhyp <- cfChange(xhyp, "age",
  x=mean(na.omit(workmom77$age))+sd(na.omit(workmom77$age)),
  xpre=mean(na.omit(workmom77$age)),
  scen=5)

xhyp <- cfName(xhyp, "Age - 1sd = 28", scen=6)
xhyp <- cfChange(xhyp, "age",
  x=mean(na.omit(workmom77$age))-sd(na.omit(workmom77$age)),
  xpre=mean(na.omit(workmom77$age)),
  scen=6)

xhyp <- cfName(xhyp, "High School Grad", scen=7)
xhyp <- cfChange(xhyp, "ed", x=12, xpre=mean(na.omit(workmom77$ed)), scen=7)

xhyp <- cfName(xhyp, "College Grad", scen=8)
xhyp <- cfChange(xhyp, "ed", x=16, xpre=mean(na.omit(workmom77$ed)), scen=8)

xhyp <- cfName(xhyp, "High Prestige Job (+1 sd)", scen=9)
xhyp <- cfChange(xhyp, "prst",
  x=mean(na.omit(workmom77$prst))+sd(na.omit(workmom77$prst)),
  xpre=mean(na.omit(workmom77$prst)),
  scen=9)

xhyp <- cfName(xhyp, "Low Prestige Job (-1 sd)", scen=10)
xhyp <- cfChange(xhyp, "prst",
  x=mean(na.omit(workmom77$prst))-sd(na.omit(workmom77$prst)),
  xpre=mean(na.omit(workmom77$prst)),
  scen=10)
```

4. Simulating predicted values and confidence intervals

```
# Simulate expected probabilities (all four categories)
oprobit.ev77 <- oprobitsimev(xhyp, simbetas, cat=4)

# Simulate first differences (all four categories)
oprobit.fd77 <- oprobitsimfd(xhyp, simbetas, cat=4)

# Simulate relative risks (all four categories)
oprobit.rr77 <- oprobitsimrr(xhyp, simbetas, cat=4)

# Plot predicted probabilities for all four categories, sorted by size
sorted <- order(oprobit.ev77$pe[,1])
scenNames <- row.names(xhyp$x)
```

4. Simulating predicted values and confidence intervals

```
trace1 <- ropeladder(x = oprobit.ev77$pe[sorted,1],
                    lower = oprobit.ev77$lower[sorted,1],
                    upper = oprobit.ev77$upper[sorted,1],
                    labels = scenNames[sorted],
                    size=0.5,
                    lex=1.5,
                    lineend="square",
                    plot=1
)

trace2 <- ropeladder(x = oprobit.ev77$pe[sorted,2],
                    lower = oprobit.ev77$lower[sorted,2],
                    upper = oprobit.ev77$upper[sorted,2],
                    size=0.5,
                    lex=1.5,
                    lineend="square",
                    plot=2
)

trace3 <- ropeladder(x = oprobit.ev77$pe[sorted,3],
                    lower = oprobit.ev77$lower[sorted,3],
                    upper = oprobit.ev77$upper[sorted,3],
                    size=0.5,
                    lex=1.5,
                    lineend="square",
                    plot=3
)
```


4. Simulating predicted values and confidence intervals

```
trace4 <- ropeladder(x = oprobit.ev77$pe[sorted,4],
                    lower = oprobit.ev77$lower[sorted,4],
                    upper = oprobit.ev77$upper[sorted,4],
                    size=0.5,
                    lex=1.5,
                    lineend="square",
                    plot=4
)

file <- "mothers4catEV"
tile(trace1, trace2, trace3, trace4,
     limits = c(0,0.5),
     gridlines = list(type="xt"),
     topaxis=list(add=TRUE, at=c(0,0.1,0.2,0.3,0.4,0.5)),
     xaxistitle=list(labels="probability"),
     topaxistitle=list(labels="probability"),
     plottitle=list(labels=c("Strongly Disagree", "Disagree",
                             "Agree", "Strongly Agree")),
     width=list(spacer=3),
     height = list(plottitle=3,xaxistitle=3.5,topaxistitle=3.5),
     output=list(outfile=file, width=12)
)
```


4. Simulating predicted values and confidence intervals

```
## Make a new rl plot, EV of Dd vs aA
trace1b <- ropeladder(x = oprobit.ev77c$pe[sorted,1],
  lower = oprobit.ev77c$lower[sorted,1],
  upper = oprobit.ev77c$upper[sorted,1],
  labels = scenNames[sorted],
  size=0.65,
  lex=1.75,
  lineend="square",
  plot=1
)

trace2b <- ropeladder(x = oprobit.ev77c$pe[sorted,2],
  lower = oprobit.ev77c$lower[sorted,2],
  upper = oprobit.ev77c$upper[sorted,2],
  size=0.65,
  lex=1.75,
  lineend="square",
  plot=2
)
```

4. Simulating predicted values and confidence intervals

```
file <- "mothers2catEV"
tile(trace1b, trace2b,
     limits = c(0.35,0.65),
     gridlines = list(type="xt"),
     xaxis=list(at=c(0.4, 0.5, 0.6)),
     topaxis=list(add=TRUE, at=c(0.4, 0.5, 0.6)),
     xaxisitle=list(labels="probability"),
     topaxisitle=list(labels="probability"),
     plottitle=list(labels=c("Disagree or Str Disagree",
                             "Agree or Str Agree")),
     width=list(spacer=3),
     height = list(plottitle=3,xaxisitle=3.5,topaxisitle=3.5),
     output=list(outfile=file, width=7)
)

## Revise traces and plot to show only "SA/A"
trace2b$entryheight <- 0.2
trace2b$plot <- 1
trace2b$labels <- trace1b$labels
file <- "mothers1catEV"
tile(trace2b,
     limits = c(0.35,0.65),
     gridlines = list(type="xt"),
     xaxis=list(at=seq(0.35, 0.65, 0.05)),
     topaxis=list(add=TRUE, at=seq(0.35, 0.65, 0.05)),
     xaxisitle=list(labels="probability agree or strongly agree, 1977"),
     topaxisitle=list(labels="probability agree or strongly agree, 1977"),
     plottitle=list(labels="\nWorking Mothers Can Be Warm to Kids\n"),
     width=list(plot=2.5),
     height = list(plottitle=3,xaxisitle=3.5,topaxisitle=3.5),
     output=list(outfile=file, width=7)
)
```

4. Simulating predicted values and confidence intervals

```
## Now estimate 1989 model
y <- workmom89$warm      # Mother can have warm feelings towards child?

x <- cbind(workmom89$male, workmom89$white, workmom89$sage,
            workmom89$ed, workmom89$prst)

# Use optim directly to get MLE
ls.result <- lm(model, data=workmom89)    # use ls estimates as starting values
stval <- c(coef(ls.result),1,2)           # initial guesses
oprobit.res89 <- optim(stval, llk.oprobit4, method="BFGS", x=x, y=y, hessian=TRUE)
pe89 <- oprobit.res89$par                 # point estimates
vc89 <- solve(oprobit.res89$hessian)      # var-cov matrix
se89 <- sqrt(diag(vc89))                  # standard errors
ll89 <- -oprobit.res89$value               # likelihood at maximum

simbetas89 <- mvnrm(sims, pe89, vc89)     # draw parameters, using MASS::mvnrm
```

4. Simulating predicted values and confidence intervals

```
# Create example counterfactuals -- for diffs
xhyp <- cfMake(model, workmom77, nscen=5)

xhyp <- cfName(xhyp, "Female (Male)", scen=1)
xhyp <- cfChange(xhyp, "male", x=0, xpre=1, scen=1)

xhyp <- cfName(xhyp, "Nonwhite (White)", scen=2)
xhyp <- cfChange(xhyp, "white", x=0, xpre=1, scen=2)

xhyp <- cfName(xhyp, "28 Year Olds (61)", scen=3)
xhyp <- cfChange(xhyp, "age",
                 x=mean(na.omit(workmom77$age))-sd(na.omit(workmom77$age)),
                 xpre=mean(na.omit(workmom77$age)),
                 scen=3)

xhyp <- cfName(xhyp, "College Grad (High School)", scen=4)
xhyp <- cfChange(xhyp, "ed", x=16, xpre=12, scen=4)

xhyp <- cfName(xhyp, "High Prestige Job (Low)", scen=5)
xhyp <- cfChange(xhyp, "prst",
                 x=mean(na.omit(workmom77$prst))+sd(na.omit(workmom77$prst)),
                 xpre=mean(na.omit(workmom77$prst)) - sd(na.omit(workmom77$prst)),
                 scen=5)
```


4. Simulating predicted values and confidence intervals

Make a new ropeladder plot, showing just change in probability of any agreement

```
sortedc <- rev(order(oprobit.fd77c$pe[,2]))
```

```
scenNames <- row.names(xhyp$x)
```

```
trace1c <- ropeladder(x = oprobit.fd77c$pe[sortedc,2],  
  lower = oprobit.fd77c$lower[sortedc,2],  
  upper = oprobit.fd77c$upper[sortedc,2],  
  labels = scenNames[sortedc],  
  sublabels="1977",  
  sublabelsyoffset=0.04,  
  col=orange,  
  size=0.65,  
  lex=1.75,  
  lineend="square",  
  plot=1  
)
```

```
trace2c <- ropeladder(x = oprobit.fd89c$pe[sortedc,2],  
  lower = oprobit.fd89c$lower[sortedc,2],  
  upper = oprobit.fd89c$upper[sortedc,2],  
  labels = scenNames[sortedc],  
  sublabels = "1989",  
  sublabelsyoffset = -0.04,  
  col=blue,  
  size=0.65,  
  lex=1.75,  
  lineend="square",  
  entryheight=0.40,  
  subentryheight=.8,  
  plot=1  
)
```


4. Simulating predicted values and confidence intervals

```
sigMark1 <- oprobit.fd77c$pe[sortedc,2]
is.na(sigMark1) <- (oprobit.fd77c$lower[sortedc,2]>0)
traceSig1 <- ropeladder(x=sigMark1,
                      col="white",
                      group=1,
                      plot=1)

sigMark2 <- oprobit.fd89c$pe[sortedc,2]
is.na(sigMark2) <- (oprobit.fd89c$lower[sortedc,2]>0)
traceSig2 <- ropeladder(x=sigMark2,
                      col="white",
                      group=2,
                      plot=1)

vertmark <- linesTile(x=c(0,0), y=c(0,1), plot=1)

file <- "mothersFD7789"
tile(trace1c, trace2c, vertmark, traceSig1, traceSig2,
     limits=c(-0.05,0.25),
     gridlines=list(type="xt"),
     topaxis=list(add=TRUE, at=seq(from=0, to=0.2, by=0.05),
                  labels=c("0%", "+5%", "+10%", "+15%", "+20%")),
     xaxis=list(at=seq(from=0, to=0.2, by=0.05), labels=c("0%", "+5%", "+10%", "+15%", "+20%")),
     xaxistitle=list(labels="difference in probability agree or strongly agree"),
     topaxistitle=list(labels="difference in probability agree or strongly agree"),
     plottitle=list(labels="\nWorking Mothers Can Be Warm to Kids\n"),
     width=list(plot=2),
     height=list(plottitle=3,xaxistitle=3.5,topaxistitle=3.5),
     output=list(outfile=file, width=6.75)
)
```